

# Retail Analytics Customer Segmentation Analysis

**Spring 2025, University of West Georgia**

**By**

**Enita Omuvwie**

# Table of contents

01

Introduction

02

Objective

03

Methodology

04

Code & Output

05

Conclusion

# Introduction

- The dataset contains **8068** entries and **9** columns with customer attributes like demographics and behaviors
- Explore customer attributes, segment based on characteristics, provide insights and inform strategic business decisions

## Description of Features:

- **Gender:** Categorical variable indicating the gender of the customer (e.g., Male, Female).
- **Ever\_Married:** Categorical variable indicating whether the customer has ever been married (e.g., Yes, No).
- **Age:** Numerical variable representing the age of the customer, measured in years.
- **Graduated:** Categorical variable indicating whether the customer has graduated from a higher education institution (e.g., Yes, No).
- **Profession:** Categorical variable representing the customer's profession (e.g., Engineer, Teacher).
- **Work\_Experience:** Numerical variable indicating the number of years of work experience the customer has.
- **Spending\_Score:** Categorical variable that rates the customer's spending behavior (e.g., Low, Medium, High).
- **Family\_Size:** Numerical variable indicating the size of the customer's family.
- **Var\_1:** Categorical variable for customer

# Objective

- The primary objective of this analysis is to perform customer segmentation using K-Means clustering.
- By identifying distinct customer groups, we aim to uncover insights that can inform marketing strategies and enhance customer engagement.

# Methodology

## Data Preprocessing

- Load in datasets (Test and Train)
- Checking for missing values
- Handling missing values

## Data Exploration

- Generate summary statistics
- Visualizing variable distributions for numerical and categorical features
- Correlation plot using heatmap

## Clustering

- Choose K-means algorithm for segmentation
- Determine optimal no. of cluster using elbow method
- Trained K-means model on scaled training data
- Assign cluster labels to both datasets

## Customer Profiling

- Create profile for each segment based on characteristics
- Analyze key features like age, income

## Result Interpretation

- Summarized findings and insights for cluster analysis
- Provide actionable recommendations based on customer segments

# Code and Output : Data Preprocessing

## Loading Libraries and Preprocessing Function to load in datasets

### Enita Omuvwie Customer Segmentation

```
[70]: # Reading in the Libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import silhouette_score
```

```
[71]: # Function to Load data
def load_data(train_path, test_path):
    train_data = pd.read_csv(train_path)
    test_data = pd.read_csv(test_path)
    return train_data, test_data
```

```
[72]: # Load datasets
train_data, test_data = load_data('Train1.csv', 'Test1.csv')

# Preprocessing
train_data = train_data.drop(['Segmentation', 'ID'], axis="columns")
test_data = test_data.drop(['ID'], axis="columns")
```

```
[73]: # Check first few rows of train data
print(train_data.head())
```

	Gender	Ever_Married	Age	Graduated	Profession	Work_Experience	\
0	Male	No	22	No	Healthcare	1.0	
1	Female	Yes	38	Yes	Engineer	NaN	
2	Female	Yes	67	Yes	Engineer	1.0	
3	Male	Yes	67	Yes	Lawyer	0.0	
4	Female	Yes	40	Yes	Entertainment	NaN	

	Spending_Score	Family_Size	Var_1
0	Low	4.0	Cat_4
1	Average	3.0	Cat_4
2	Low	1.0	Cat_6
3	High	2.0	Cat_6
4	High	6.0	Cat_6

# Code and Output : Data Preprocessing

Checking the data overview and descriptive statistics

```
[74]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8068 entries, 0 to 8067
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Gender       8068 non-null    object  
 1   Ever_Married 7928 non-null    object  
 2   Age          8068 non-null    int64  
 3   Graduated    7990 non-null    object  
 4   Profession   7944 non-null    object  
 5   Work_Experience 7239 non-null  float64 
 6   Spending_Score 8068 non-null  object  
 7   Family_Size   7733 non-null    float64 
 8   Var_1         7992 non-null    object  
dtypes: float64(2), int64(1), object(6)
memory usage: 567.4+ KB
```

```
[75]: train_data.describe().T
```

```
[75]:
```

	count	mean	std	min	25%	50%	75%	max
<b>Age</b>	8068.0	43.466906	16.711696	18.0	30.0	40.0	53.0	89.0
<b>Work_Experience</b>	7239.0	2.641663	3.406763	0.0	0.0	1.0	4.0	14.0
<b>Family_Size</b>	7733.0	2.850123	1.531413	1.0	2.0	3.0	4.0	9.0

```
[76]: train_data.describe(include='object').T
```

```
[76]:
```

	count	unique	top	freq
<b>Gender</b>	8068	2	Male	4417
<b>Ever_Married</b>	7928	2	Yes	4643
<b>Graduated</b>	7990	2	Yes	4968
<b>Profession</b>	7944	9	Artist	2516
<b>Spending_Score</b>	8068	3	Low	4878
<b>Var_1</b>	7992	7	Cat_6	5238

# Code and Output : Data Preprocessing

Checking the data overview and descriptive statistics

```
[77]: test_data.head()
[77]:
   Gender Ever_Married  Age  Graduated Profession Work_Experience  Spending_Score  Family_Size  Var_1
0  Female        Yes  36       Yes  Engineer           0.0        Low      1.0  Cat_6
1    Male        Yes  37       Yes  Healthcare          8.0  Average      4.0  Cat_6
2  Female        Yes  69       No     NaN            0.0        Low      1.0  Cat_6
3    Male        Yes  59       No  Executive         11.0        High      2.0  Cat_6
4  Female        No  19       No  Marketing          NaN        Low      4.0  Cat_6

[78]: test_data.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2627 entries, 0 to 2626
Data columns (total 9 columns):
 #   column          non-null Count  Dtype  
 --  --  
 0   Gender          2627 non-null   object  
 1   Ever_Married    2577 non-null   object  
 2   Age              2627 non-null   int64  
 3   Graduated        2603 non-null   object  
 4   Profession      2589 non-null   object  
 5   Work_Experience  2358 non-null   float64 
 6   Spending_Score  2514 non-null   float64 
 7   Family_Size      2514 non-null   float64 
 8   Var_1             2595 non-null   object  
dtypes: float64(2), int64(1), object(6)
memory usage: 184.8+ KB

[79]: test_data.describe()
[79]:
      Age  Work_Experience  Family_Size
count  2627.000000  2358.000000  2514.000000
mean   43.649791  2.552587  2.825378
std    16.967015  3.341094  1.551906
min    18.000000  0.000000  1.000000
25%   30.000000  0.000000  2.000000
50%   41.000000  1.000000  2.000000
75%   53.000000  4.000000  4.000000
max    89.000000  14.000000  9.000000

[80]: test_data.describe(include="object").T
[80]:
      count  unique  top  freq
Gender  2627      2  Male  1424
Ever_Married  2577      2  Yes  1520
Graduated  2603      2  Yes  1602
Profession  2589      9  Artist  802
Spending_Score  2627      3  Low  1616
Var_1    2595      7  Cat_6  1672
```

# Code and Output : Data Preprocessing

## Checking for missing values

```
# Function to plot missing values
def plot_missing_values(data, dataset_name="Dataset"):
    """
    Plots the percentage of missing values for each column in the dataset.

    Parameters:
    - data: Pandas DataFrame containing the dataset
    - dataset_name: String, name of the dataset to display in the title
    """
    missing_data = data.isnull().sum()
    missing_percent = (missing_data[missing_data > 0] / data.shape[0]) * 100
    missing_percent.sort_values(ascending=True, inplace=True)

    if missing_percent.empty:
        print(f"No missing values in {dataset_name}!")
        return

    fig, ax = plt.subplots(figsize=(15, 4))
    ax.barh(missing_percent.index, missing_percent, color='blue')

    for i, (value, name) in enumerate(zip(missing_percent, missing_percent.index)):
        ax.text(value + 0.5, i, f"{value:.2f}%", ha='left', va='center', fontweight='bold', fontsize=12, color='black')

    ax.set_xlim([0, min(100, max(missing_percent) + 5)]) # Dynamically adjust x-axis Limit
    plt.title(f"Percentage of Missing Values - {dataset_name}", fontweight='bold', fontsize=16)
    plt.xlabel('Percentages (%)', fontsize=12)
    plt.show()

# Plot missing values
plot_missing_values(train_data, "Train Dataset")
plot_missing_values(test_data, "Test Dataset")
```

# Code and Output : Data Preprocessing

## Checking for missing values

```
# Function to plot missing values
def plot_missing_values(data, dataset_name="Dataset"):
    """
    Plots the percentage of missing values for each column in the dataset.

    Parameters:
    - data: Pandas DataFrame containing the dataset
    - dataset_name: String, name of the dataset to display in the title
    """
    missing_data = data.isnull().sum()
    missing_percent = (missing_data[missing_data > 0] / data.shape[0]) * 100
    missing_percent.sort_values(ascending=True, inplace=True)

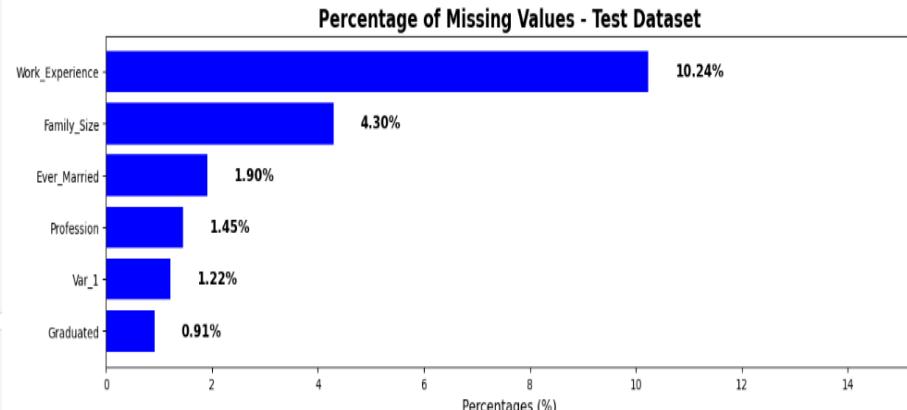
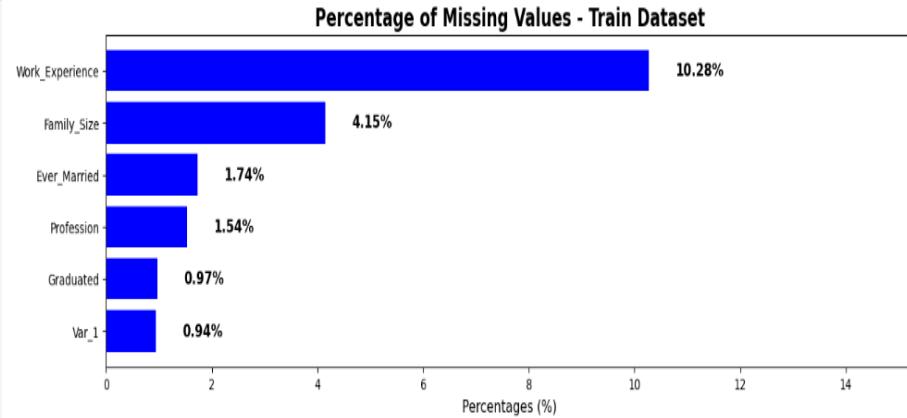
    if missing_percent.empty:
        print(f"No missing values in {dataset_name}!")
        return

    fig, ax = plt.subplots(figsize=(15, 4))
    ax.barh(missing_percent.index, missing_percent, color='blue')

    for i, (value, name) in enumerate(zip(missing_percent, missing_percent.index)):
        ax.text(value + 0.5, i, f'{value:.2f}%', ha='left', va='center', fontweight='bold', fontsize=12, color='black')

    ax.set_xlim([0, min(100, max(missing_percent) + 5)]) # Dynamically adjust x-axis limit
    plt.title(f"Percentage of Missing Values - {dataset_name}", fontweight='bold', fontsize=16)
    plt.xlabel('Percentages (%)', fontsize=12)
    plt.show()

# Plot missing values
plot_missing_values(train_data, "Train Dataset")
plot_missing_values(test_data, "Test Dataset")
```



# Code and Output : Data Preprocessing

## Handling missing values

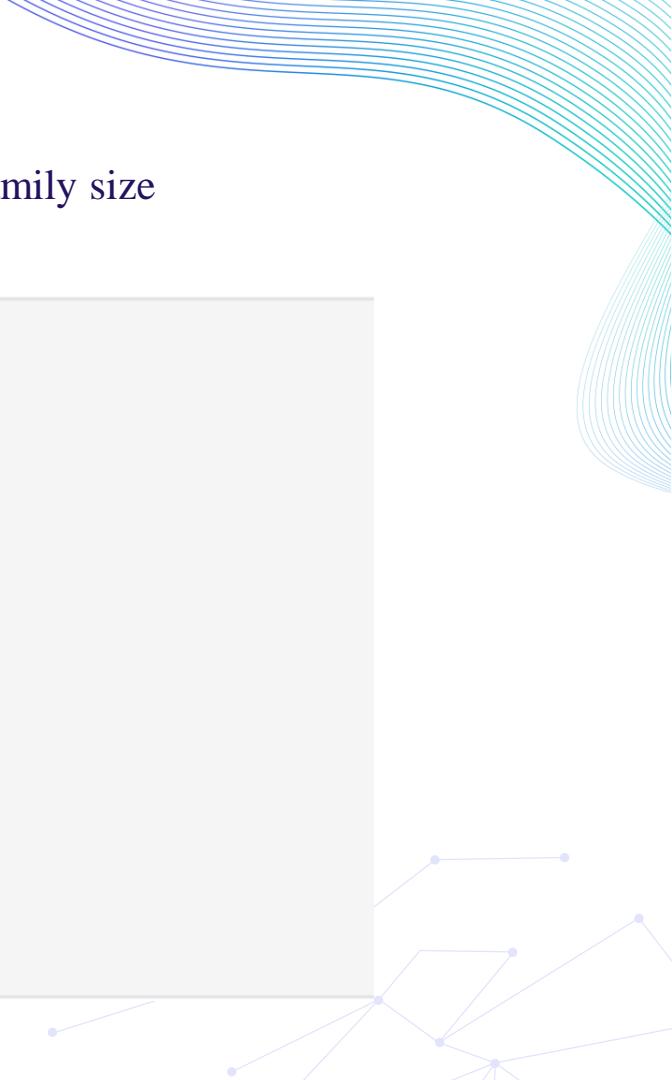
```
|: # Function to handle missing values
| def handle_missing_values(data, categorical_cols, numerical_cols):
|     for col in categorical_cols:
|         data[col].fillna(data[col].mode()[0], inplace=True)
|     for col in numerical_cols:
|         data[col].fillna(data[col].median(), inplace=True)
|     return data
|
# Handle missing values
categorical_cols = ['Ever_Married', 'Graduated', 'Profession', 'Var_1']
numerical_cols = ['Work_Experience', 'Family_Size']
train_data = handle_missing_values(train_data, categorical_cols, numerical_cols)
test_data = handle_missing_values(test_data, categorical_cols, numerical_cols)
```

# Code and Output : Data Exploration

Visualizing distribution of work experience among customers and family size

```
# Visualize distributions of numerical features
plt.figure(figsize=(12, 6))
for i, col in enumerate(numerical_cols):
    plt.subplot(1, 2, i + 1)
    sns.histplot(train_data[col], bins=20, kde=True)
    plt.title(f'Distribution of {col}')
plt.tight_layout()
plt.show()

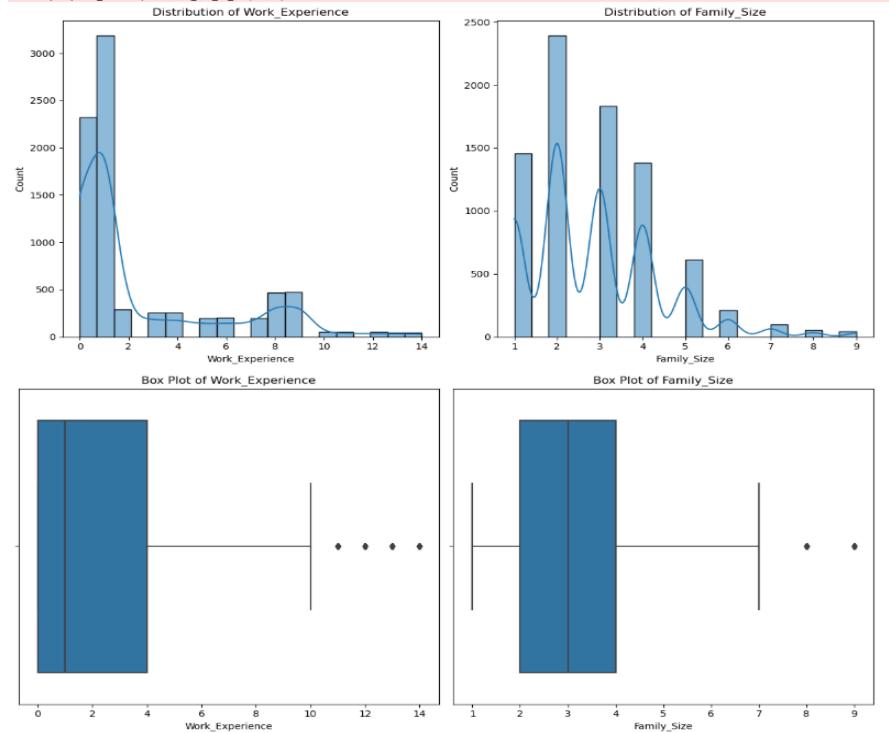
# Box plots for numerical features
plt.figure(figsize=(12, 6))
for i, col in enumerate(numerical_cols):
    plt.subplot(1, 2, i + 1)
    sns.boxplot(x=train_data[col])
    plt.title(f'Box Plot of {col}')
plt.tight_layout()
plt.show()
```



# Code and Output : Data Exploration

Visualizing distribution of work experience among customers, showing right skewed pattern for customers less than 5 years of experience. Box plot shows a central tendency for work experience.

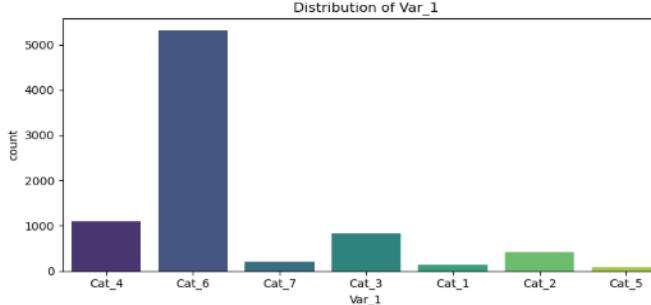
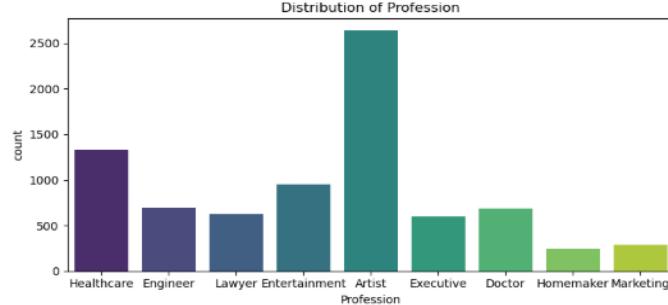
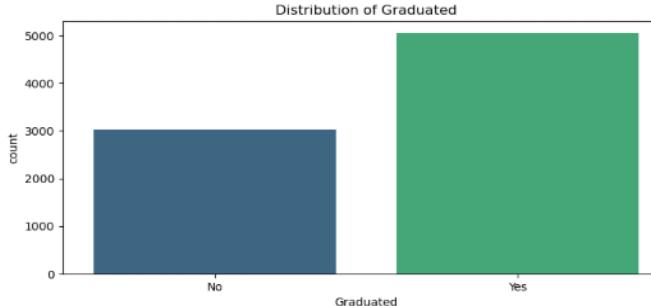
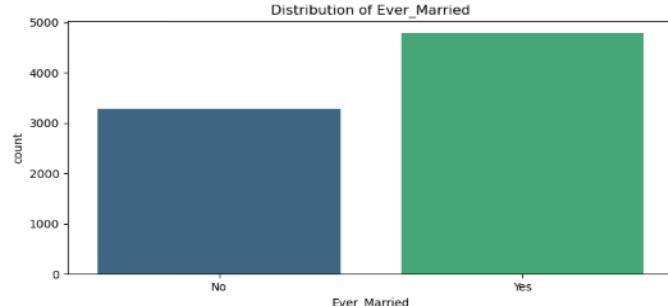
Family size indicates customers have between 2-4 members and the box plot shows the median of the data size



# Code and Output : Data Exploration

Visualizing categorical variables showing higher proportion of counts in different variable values

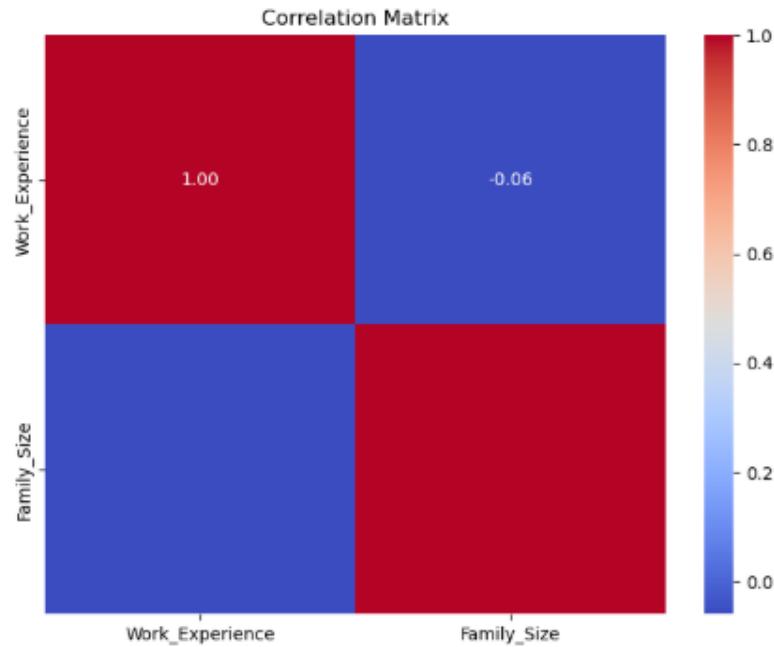
```
# Visualize distributions of categorical features
plt.figure(figsize=(16, 8))
for i, col in enumerate(categorical_cols):
    plt.subplot(2, 2, i + 1)
    sns.countplot(data=train_data, x=col, palette='viridis')
    plt.title(f'Distribution of {col}')
plt.tight_layout()
plt.show()
```



# Code and Output : Data Exploration

Correlation Analysis: Indicating weak correlation between Family Size and Work Experience

```
# Correlation Analysis
correlation_matrix = train_data[numerical_cols].corr()
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix')
plt.show()
```



# Code and Output : Data Exploration

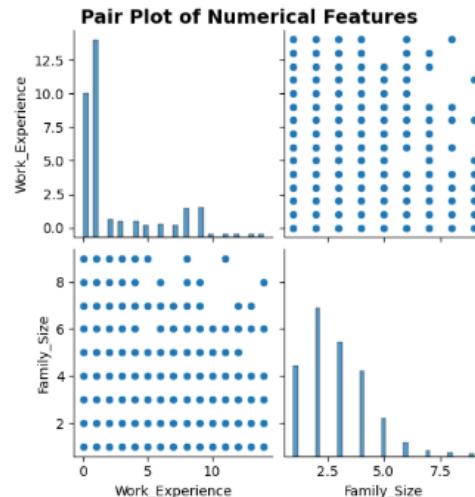
Pair Plot Analysis showing the relationship between numerical features allowing the identification of potential correlations and distributions

```
]: # Pair plot
plt.figure(figsize=(10, 6))
sns.pairplot(train_data[numerical_cols]) # Pair plot

# Set title correctly for the entire figure
plt.suptitle('Pair Plot of Numerical Features', fontweight='bold', fontsize=14, y=1.02)

plt.show()

/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Instead, set the low_memory parameter.
  with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/envs/anaconda-2024.02-py310/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Instead, set the low_memory parameter.
  with pd.option_context('mode.use_inf_as_na', True):
<Figure size 1000x600 with 0 Axes>
```



# Code and Output : Clustering

**Label Encoding:** Converted variables into numerical format using one-hot encoding

**Feature Scaling-Standardization:** Standardized the dataset by transforming features to have a mean of 0 and standard deviation of 1

## Label Encoding

```
[85]: # Function to encode categorical variables
def encode_data(data):
    return pd.get_dummies(data, drop_first=True)
```

```
[86]: # Encode categorical variables
train_encoded = encode_data(train_data)
test_encoded = encode_data(test_data)
```

## Feature Scaling - Standardization

```
[87]: # Function to scale data
def scale_data(data):
    scaler = StandardScaler()
    return scaler.fit_transform(data)
```

```
[88]: # Scale data
scaled_train = scale_data(train_encoded)
scaled_test = scale_data(test_encoded)
```

# Code and Output : Clustering

**K-Means Clustering Function:** This function trains the model and returns predicted cluster label

**Clustering Execution:** The model trained on 3 clusters on the scaled training data and resulting cluster labels added

**Cluster Prediction for Test Data:** The trained model is applied to the scaled test to predict cluster labels added to test dataset

```
# Function to perform K-Means clustering
def perform_kmeans(data, n_clusters):
    kmeans_model = KMeans(n_clusters=n_clusters, random_state=42)
    return kmeans_model.fit_predict(data)

# Perform K-Means clustering
n_clusters = 3
train_clusters = perform_kmeans(scaled_train, n_clusters)

# Add cluster Labels to the dataset
train_encoded['Cluster'] = train_clusters

# Predict clusters for test data
test_clusters = perform_kmeans(scaled_test, n_clusters)
test_encoded['Cluster'] = test_clusters
```



# Code and Output : Clustering

## Cluster Distribution Visualization: Body of Code

```
# Visualize cluster distribution
sns.countplot(x=train_encoded['Cluster'], palette='viridis')
plt.title("Cluster Distribution in Train Data")
plt.show()

# Summary statistics per cluster
cluster_summary = train_encoded.groupby('Cluster').mean()
print(cluster_summary)

# Display the count of customers in each cluster
cluster_counts = train_encoded['Cluster'].value_counts()
print(cluster_counts)

# Elbow method to determine optimal k
ssd = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    kmeans.fit(scaled_train)
    ssd.append(kmeans.inertia_)

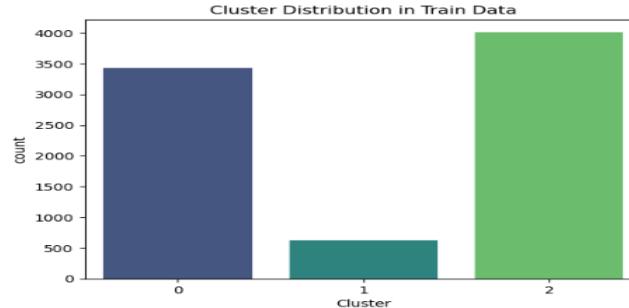
# Plot Elbow Method
plt.figure(figsize=(8, 5))
plt.plot(range(1, 11), ssd, marker='o', linestyle='--')
plt.xlabel('Number of Clusters')
plt.ylabel('SSD')
plt.title('Elbow Method to Determine Optimal k')
plt.show()
```



# Code and Output : Clustering

**Summary Statistics per Cluster:** Shows the statistics for each cluster providing insights into the average customer characteristics within each segment

**Customer Count in Each Cluster:** The count is displayed offering a view of the size of each segment



Cluster	Age	Work_Experience	Family_Size	Gender_Male
0	31.715701	2.881154	3.001165	0.481794
1	75.464516	1.191935	1.975806	0.509677
2	48.573599	2.321793	2.868493	0.609465

Cluster	Ever_Married_Yes	Graduated_Yes	Profession_Doctor
0	0.655928	0.503059	0.116867
1	0.938710	0.633871	0.000000
2	0.998506	0.728767	0.671482

Cluster	Profession_Engineer	Profession_Entertainment	Profession_Executive
0	0.082726	0.105738	0.009904
1	0.000000	0.000000	0.000000
2	0.103362	0.145953	0.148722

Cluster	...	Profession_Lawyer	Profession_Marketing	Spending_Score_High
0	...	0.000874	0.061462	0.000000
1	...	1.000000	0.000000	0.522581
2	...	0.000000	0.020174	0.222167

Cluster	Spending_Score_Low	Var_1_Cat_2	Var_1_Cat_3	Var_1_Cat_4
0	0.099835	0.073405	0.113603	0.147684
1	0.448387	0.012903	0.045161	0.058065
2	0.291656	0.048349	0.100623	0.135990

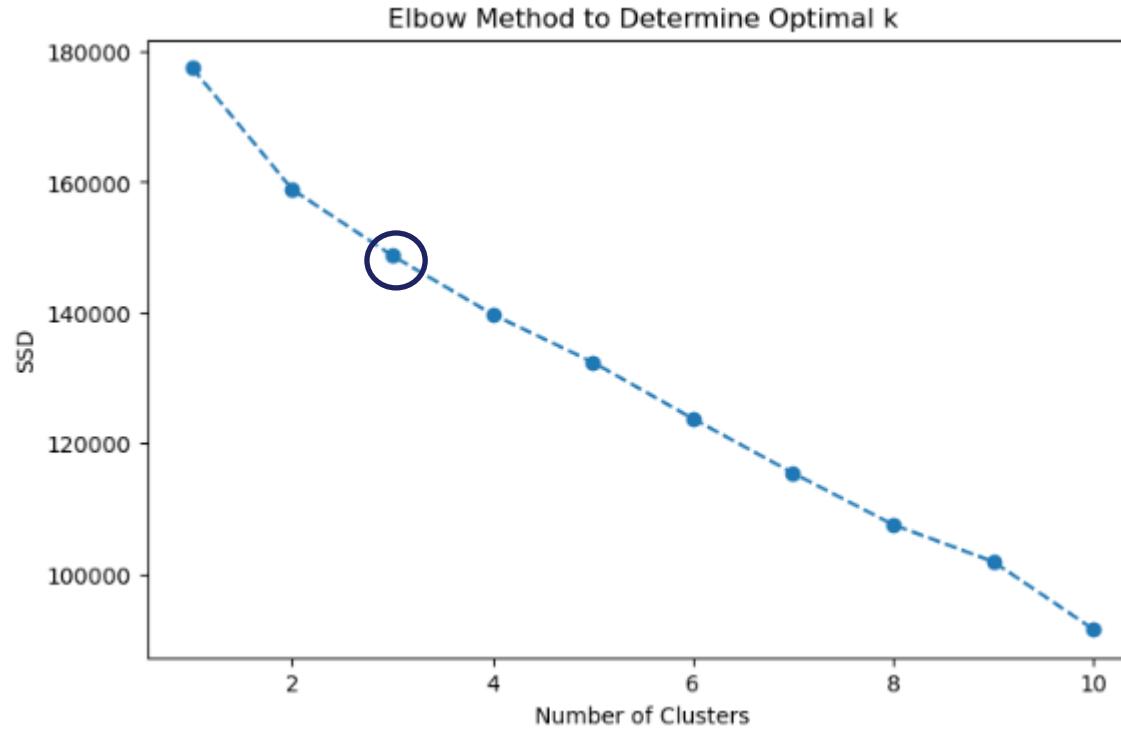
  

Cluster	Var_1_Cat_5	Var_1_Cat_6	Var_1_Cat_7
0	0.014856	0.683554	0.829428
1	0.004339	0.862903	0.000000
2	0.007721	0.674222	0.025405



# Code and Output : Clustering

**Elbow Method for Optimal K:** The elbow method is implemented to determine optimal number of clusters using sum of squared distances (SSD) for different k values



# Code and Output : Clustering

**Customer Profiles:** Containing a detail of each cluster profile with a summary mean statistics, providing more insights to customer characteristics in each segment

```
[]: # Create a detailed profile for each cluster
for cluster in range(cluster_summary.shape[0]):
    print(f"Profile for Cluster {cluster}:")
    print(f"Average Age: {cluster_summary.loc[cluster, 'Age']:.2f}")
    print(f"Average Work Experience: {cluster_summary.loc[cluster, 'Work_Experience']:.2f} years")
    print(f"Average Family Size: {cluster_summary.loc[cluster, 'Family_Size']:.2f}")
    print(f"Proportion Ever Married: {cluster_counts[cluster] / train_encoded.shape[0]:.2%}")
    print("\n")
```

```
Profile for Cluster 0:
Average Age: 31.72
Average Work Experience: 2.88 years
Average Family Size: 3.00
Proportion Ever Married: 42.55%
```

```
Profile for Cluster 1:
Average Age: 75.46
Average Work Experience: 1.19 years
Average Family Size: 1.98
Proportion Ever Married: 7.68%
```

```
Profile for Cluster 2:
Average Age: 48.57
Average Work Experience: 2.32 years
Average Family Size: 2.87
Proportion Ever Married: 49.76%
```



# Code and Output : Clustering

**Silhouette Score Calculation:** Used to assess the quality of clustering with 0.13 indicating clusters are not well separated

**T-SNE:** Shows the cluster in 2D space showing the distribution of clusters

```
# Calculate the silhouette score
silhouette_avg = silhouette_score(scaled_train, train_clusters)
print(f"Silhouette Score: {silhouette_avg:.2f}")

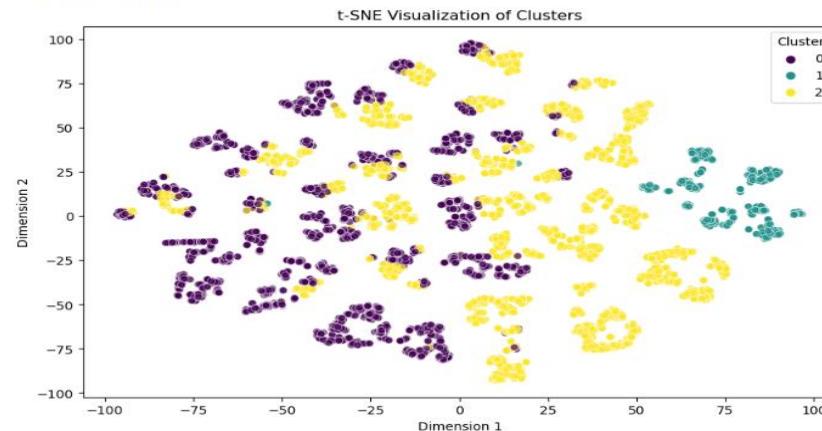
# Optional: Visualize the silhouette scores for each sample
from sklearn.manifold import TSNE

# Reduce dimensions for visualization
tsne = TSNE(n_components=2, random_state=42)
tsne_results = tsne.fit_transform(scaled_train)

# Create a DataFrame for visualization
tsne_df = pd.DataFrame(tsne_results, columns=['Dimension 1', 'Dimension 2'])
tsne_df['Cluster'] = train_clusters

# Plot the t-SNE results
plt.figure(figsize=(10, 6))
sns.scatterplot(data=tsne_df, x='Dimension 1', y='Dimension 2', hue='cluster', palette='viridis', alpha=0.7)
plt.title("t-SNE Visualization of Clusters")
plt.show()
```

Silhouette Score: 0.13



# Conclusion

- **Optimal Clusters:** The Elbow method suggests 3 clusters for effective customer segmentation
- **Distinct Profiles:** Significant differences was identified in demographics and behaviors among clusters
- **Silhouette Score:** a score of 0.13 indicates the clusters may not be properly separated
- **Further analysis:** other algorithms can be used to improve the model
- **Recommendations:** Insights can be used for marketing strategies and targeted customers-based advertising based on the results



# Thanks!

**Do you have any questions?**  
[eomuvwi1@my.westga.edu](mailto:eomuvwi1@my.westga.edu)